

# Data Wrangling with dplyr

Stats 102A

Miles Chen

Department of Statistics

Week 4 Monday



# grumpy professor chat

It is your responsibility to read through your entire HW output file before you submit it. Make sure all requested output is visible.

I've received a few regrade requests for submissions that look like this, which unfortunately will not earn points in a regrade.

```
month_names <- read.delim("month_names.txt", encoding="UTF-8", row.names=1)

## Warning in file(file, "rt"): cannot open file 'month_names.txt': No such file or
## directory

## Error in file(file, "rt"): cannot open the connection

x <- factor(c("March", "March", "February", "June"))
month_convert(x, "English", "Spanish")

## Warning in file(file, "rt"): cannot open file 'month_names.txt': No such file or
## directory

## Error in file(file, "rt"): cannot open the connection
```

# A silly mistake can make a big difference

It is not fun losing points because of a silly mistake. A silly mistake on the HW might cost you a percentage point or two in your final course grade. This is a relatively small price to pay and (in my opinion) worth the consequence if it motivates someone to pay more attention to detail.

Many job applications will say they need someone “detail oriented.” This means that silly mistakes can have large consequences and they need someone who double or triple checks their own work before submission.

Imagine the consequences of:

- Submitting the wrong version of your personal statement on a school application
- CCing the wrong person on an email with confidential information
- Entering the wrong digit of a bank account number on transfer order
- An extra 0 in the dosage of a medication

I want you to achieve your dreams and potential. I don't want silly mistakes to hold you back.

## Section 1

dplyr

dplyr is a core part of the tidyverse.

You can load the library with `library(dplyr)` or by loading all of the tidyverse with `library(tidyverse)`

When working with data you must:

- Figure out what you want to do.
- Describe those tasks in the form of a computer program.
- Execute the program.

The dplyr package makes these steps fast and easy:

- By constraining your options, it helps you think about your data manipulation challenges.
- It provides simple “verbs”, functions that correspond to the most common data manipulation tasks, to help you translate your thoughts into code.

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

- `select()` picks variables based on their names.
- `filter()` picks cases based on their values.
- `mutate()` adds new variables that are functions of existing variables.
- `arrange()` changes the ordering of the rows.
- `summarise()` reduces multiple values down to a single summary.

These all combine naturally with `group_by()` which allows you to perform any operation “by group.”





## Section 2

### dplyr examples

# The starwars data set

The Star Wars data set is included with `dplyr`. It contains information about various Star Wars characters from the first 7 Star Wars movies.

```
starwars
```

```
## # A tibble: 87 x 14
##   name height mass hair_color skin_color eye_color birth_year sex gender homeworld species
##   <chr> <int> <dbl> <chr> <chr> <chr> <dbl> <chr> <chr> <chr> <chr>
## 1 Luke~ 172 77 blond fair blue 19 male mascu~ Tatooine Human
## 2 C-3P0 167 75 <NA> gold yellow 112 none mascu~ Tatooine Droid
## 3 R2-D2 96 32 <NA> white, bl~ red 33 none mascu~ Naboo Droid
## 4 Dart~ 202 136 none white yellow 41.9 male mascu~ Tatooine Human
## 5 Leia~ 150 49 brown light brown 19 fema~ femin~ Alderaan Human
## 6 Owen~ 178 120 brown, gr~ light blue 52 male mascu~ Tatooine Human
## 7 Beru~ 165 75 brown light blue 47 fema~ femin~ Tatooine Human
## 8 R5-D4 97 32 <NA> white, red red NA none mascu~ Tatooine Droid
## 9 Bigg~ 183 84 black light brown 24 male mascu~ Tatooine Human
## 10 Obi~ 182 77 auburn, w~ fair blue-gray 57 male mascu~ Stewjon Human
## # ... with 77 more rows, and 3 more variables: films <list>, vehicles <list>, starships <list>
```

## Select columns with `select()`

When using `select()`, you do not need to put quotes around the column names if there are no spaces in the names.

```
select(starwars, name, homeworld, species, films)
```

```
## # A tibble: 87 x 4
##   name                homeworld species films
##   <chr>              <chr>      <chr>  <list>
## 1 Luke Skywalker    Tatooine  Human  <chr [5]>
## 2 C-3P0             Tatooine  Droid  <chr [6]>
## 3 R2-D2             Naboo     Droid  <chr [7]>
## 4 Darth Vader       Tatooine  Human  <chr [4]>
## 5 Leia Organa       Alderaan Human  <chr [5]>
## 6 Owen Lars         Tatooine  Human  <chr [3]>
## 7 Beru Whitesun lars Tatooine  Human  <chr [3]>
## 8 R5-D4             Tatooine  Droid  <chr [1]>
## 9 Biggs Darklighter Tatooine  Human  <chr [1]>
## 10 Obi-Wan Kenobi    Stewjon   Human  <chr [6]>
## # ... with 77 more rows
```

# Using the pipe

The pipe `%>%` takes the result of what is in front of the pipe and inserts it as the first argument in the function that comes after the pipe. `x %>% f(y)` turns into `f(x, y)` so the result from one step is then “piped” into the next step.

```
# select(starwars, name, homeworld, species, films) is exactly equivalent to  
starwars %>% select(name, homeworld, species, films)
```

```
## # A tibble: 87 x 4  
##   name          homeworld species films  
##   <chr>         <chr>      <chr>  <list>  
## 1 Luke Skywalker Tatooine   Human  <chr [5]>  
## 2 C-3PO         Tatooine   Droid  <chr [6]>  
## 3 R2-D2         Naboo      Droid  <chr [7]>  
## 4 Darth Vader   Tatooine   Human  <chr [4]>  
## 5 Leia Organa   Alderaan  Human  <chr [5]>  
## 6 Owen Lars     Tatooine   Human  <chr [3]>  
## 7 Beru Whitesun lars Tatooine   Human  <chr [3]>  
## 8 R5-D4         Tatooine   Droid  <chr [1]>  
## 9 Biggs Darklighter Tatooine   Human  <chr [1]>  
## 10 Obi-Wan Kenobi Stewjon    Human  <chr [6]>  
## # ... with 77 more rows
```

# Shortcut to insert the pipe

Shortcut to insert the pipe:

**CTRL(CMD) + SHIFT + M**

# Select columns with `select()`

- Use a negative sign to deselect columns

```
starwars %>%  
  select( -name, -eye_color, -birth_year) %>%  
  head(3)
```

```
## # A tibble: 3 x 11  
##   height mass hair_color skin_color sex gender homeworld species films vehicles starships  
##   <int> <dbl> <chr>      <chr>      <chr> <chr> <chr> <chr> <list> <list> <list>  
## 1    172    77 blond      fair      male masculine Tatooine Human <chr [5]> <chr [2~ <chr [2]>  
## 2    167    75 <NA>      gold      none masculine Tatooine Droid <chr [6]> <chr [0~ <chr [0]>  
## 3     96    32 <NA>      white, blue none masculine Naboo Droid <chr [7]> <chr [0~ <chr [0]>
```

# select() example

- Use colon notation to select a range of columns

```
starwars %>%  
  select(name:eye_color) %>%  
  head(3)
```

```
## # A tibble: 3 x 6  
##   name          height  mass hair_color skin_color eye_color  
##   <chr>         <int> <dbl> <chr>      <chr>      <chr>  
## 1 Luke Skywalker   172    77 blond      fair        blue  
## 2 C-3PO             167    75 <NA>       gold        yellow  
## 3 R2-D2              96    32 <NA>       white, blue red
```

# Special selection function

`dplyr` has special selection functions. See `?tidyselect::select_helpers`

- `contains()` Select columns that contain a character string
- `starts_with()` Select columns that start with a character string
- `ends_with()` Select columns that end with a string
- `matches()` Select columns that match a regular expression
- `everything()` Select all columns
- `num_range()` Select columns named something like `x1`, `x2`, `x3`, `x4`, `x5`
- `one_of(name_vector)` Select columns where the names are stored in a vector



# Selection function examples

```
starwars %>%  
  select(name, ends_with("color")) %>% # selects name and columns ending with color  
  head(3)
```

```
## # A tibble: 3 x 4  
##   name          hair_color skin_color eye_color  
##   <chr>         <chr>      <chr>    <chr>  
## 1 Luke Skywalker blond      fair      blue  
## 2 C-3P0         <NA>      gold      yellow  
## 3 R2-D2         <NA>      white, blue red
```

```
# selects name column and columns that match the regex, which says ends with "s"  
starwars %>%  
  select(name, matches("s$")) %>%  
  head(3)
```

```
## # A tibble: 3 x 6  
##   name          mass species films      vehicles starships  
##   <chr>         <dbl> <chr>  <list>   <list>    <list>  
## 1 Luke Skywalker    77 Human <chr [5]> <chr [2]> <chr [2]>  
## 2 C-3P0              75 Droid <chr [6]> <chr [0]> <chr [0]>  
## 3 R2-D2              32 Droid <chr [7]> <chr [0]> <chr [0]>
```

# Selecting with a variable

You can also select with a vector of names. To accomplish this, use the functions `all_of()` or `any_of()`

```
vars <- c("name", "mass", "height")
starwars %>% select(all_of(vars))
```

```
## # A tibble: 87 x 3
##   name          mass height
##   <chr>      <dbl> <int>
## 1 Luke Skywalker    77    172
## 2 C-3P0             75    167
## 3 R2-D2             32     96
## 4 Darth Vader      136    202
## 5 Leia Organa       49    150
## 6 Owen Lars        120    178
## 7 Beru Whitesun lars  75    165
## 8 R5-D4             32     97
## 9 Biggs Darklighter  84    183
## 10 Obi-Wan Kenobi    77    182
## # ... with 77 more rows
```

## Filter rows with filter()

With `filter()` you specify conditions to filter the rows in the data. Filter can use any condition that can be expressed as a logical vector with length equal to the number of rows.

```
starwars %>%  
  filter(name == "R2-D2")
```

```
## # A tibble: 1 x 14  
##   name height mass hair_color skin_color eye_color birth_year sex gender homeworld species films  
##   <chr>  <int> <dbl> <chr>      <chr>      <chr>          <dbl> <chr> <chr>  <chr>      <chr>  <lis>  
## 1 R2-D2    96    32 <NA>      white, bl~ red          33 none  mascu~ Naboo      Droid  <chr~  
## # ... with 2 more variables: vehicles <list>, starships <list>
```

# filter() examples

Multiple conditions can be applied. Using the comma is equivalent to using &

```
starwars %>%  
  filter(species %in% c("Human", "Droid"), height < 175)
```

```
## # A tibble: 16 x 14  
##   name height mass hair_color skin_color eye_color birth_year sex gender homeworld species  
##   <chr> <int> <dbl> <chr> <chr> <chr> <dbl> <chr> <chr> <chr> <chr>  
## 1 Luke~ 172 77 blond fair blue 19 male mascu~ Tatooine Human  
## 2 C-3P0 167 75 <NA> gold yellow 112 none mascu~ Tatooine Droid  
## 3 R2-D2 96 32 <NA> white, bl~ red 33 none mascu~ Naboo Droid  
## 4 Leia~ 150 49 brown light brown 19 fema~ femin~ Alderaan Human  
## 5 Beru~ 165 75 brown light blue 47 fema~ femin~ Tatooine Human  
## 6 R5-D4 97 32 <NA> white, red red NA none mascu~ Tatooine Droid  
## 7 Wedg~ 170 77 brown fair hazel 21 male mascu~ Corellia Human  
## 8 Palp~ 170 75 grey pale yellow 82 male mascu~ Naboo Human  
## 9 Mon ~ 150 NA auburn fair blue 48 fema~ femin~ Chandrila Human  
## 10 Fini~ 170 NA blond fair blue 91 male mascu~ Coruscant Human  
## 11 Shmi~ 163 NA black fair brown 72 fema~ femin~ Tatooine Human  
## 12 Cordé 157 NA brown light brown NA fema~ femin~ Naboo Human  
## 13 Dormé 165 NA brown light brown NA fema~ femin~ Naboo Human  
## 14 Joca~ 167 NA white fair blue NA fema~ femin~ Coruscant Human  
## 15 R4-P~ 96 NA none silver, r~ red, blue NA none femin~ <NA> Droid
```

## filter() is very powerful with regular expressions

We'll learn regular expressions in the next lecture. `str_detect()` returns a logical vector.

```
starwars %>%  
  filter(str_detect(name, "^F")) # the name starts with F
```

```
## # A tibble: 2 x 14  
##   name height mass hair_color skin_color eye_color birth_year sex gender homeworld species films  
##   <chr>  <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  <chr>      <chr>  <lis>  
## 1 Fini~   170    NA blond      fair       blue        91 male mascul~ Coruscant Human  <chr~  
## 2 Finn    NA    NA black      dark       dark        NA male mascul~ <NA>      Human  <chr~  
## # ... with 2 more variables: vehicles <list>, starships <list>
```

# The dplyr functions can be piped into each other

- use | for 'OR'

```
starwars %>%  
  filter(hair_color == "none" | eye_color == "black") %>%  
  select(name, species, homeworld, hair_color, eye_color)
```

```
## # A tibble: 38 x 5
```

	name	species	homeworld	hair_color	eye_color
	<chr>	<chr>	<chr>	<chr>	<chr>
## 1	Darth Vader	Human	Tatooine	none	yellow
## 2	Greedo	Rodian	Rodia	<NA>	black
## 3	IG-88	Droid	<NA>	none	red
## 4	Bossk	Trandosha	Trandosha	none	red
## 5	Lobot	Human	Bespin	none	blue
## 6	Ackbar	Mon Calamari	Mon Cala	none	orange
## 7	Nien Nunb	Sullustan	Sullust	none	black
## 8	Nute Gunray	Neimodian	Cato Neimoidia	none	red
## 9	Jar Jar Binks	Gungan	Naboo	none	orange
## 10	Roos Tarpals	Gungan	Naboo	none	orange
## #	... with 28 more rows				

# Sort rows with arrange()

If you want to put things in descending order, wrap the variable name with desc()

```
starwars %>%  
  select(name, birth_year, height, mass) %>%  
  arrange(desc(birth_year), mass)
```

```
## # A tibble: 87 x 4  
##   name                birth_year height  mass  
##   <chr>              <dbl>   <int> <dbl>  
## 1 Yoda                896     66    17  
## 2 Jabba Desilijic Tiure    600    175  1358  
## 3 Chewbacca           200    228   112  
## 4 C-3PO               112    167    75  
## 5 Dooku               102    193    80  
## 6 Ki-Adi-Mundi         92     198    82  
## 7 Qui-Gon Jinn         92     193    89  
## 8 Finis Valorum        91     170    NA  
## 9 Palpatine            82     170    75  
## 10 Cliegg Lars          82     183    NA  
## # ... with 77 more rows
```

# Select rows based on their position with `slice()`

`slice()` lets you select rows based on their locations. The following selects rows 5 through 10

```
starwars %>% slice(5:10)
```

```
## # A tibble: 6 x 14
##   name height mass hair_color skin_color eye_color birth_year sex gender homeworld species films
##   <chr>  <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>  <chr>      <chr>  <lis>
## 1 Leia~   150    49 brown      light      brown         19 fema~ femin~ Alderaan Human  <chr~
## 2 Owen~   178   120 brown, gr~ light      blue          52 male  mascu~ Tatooine Human  <chr~
## 3 Beru~   165    75 brown      light      blue          47 fema~ femin~ Tatooine Human  <chr~
## 4 R5-D4    97    32 <NA>      white, red red          NA none  mascu~ Tatooine Droid   <chr~
## 5 Bigg~   183    84 black      light      brown         24 male  mascu~ Tatooine Human  <chr~
## 6 Obi-~   182    77 auburn, w~ fair      blue-gray     57 male  mascu~ Stewjon Human  <chr~
## # ... with 2 more variables: vehicles <list>, starships <list>
```



# slice\_sample()

`slice_sample()` lets you randomly select rows which can be useful to get a peek at portions of the entire tibble rather than just the head

```
starwars %>% slice_sample(n = 5)
```

```
## # A tibble: 5 x 14
```

```
##   name height mass hair_color skin_color eye_color birth_year sex gender homeworld species films
##   <chr> <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr> <chr>      <chr> <lis>
## 1 San ~    191   NA none      grey      gold        NA male mascu~ Muunilin~ Muun   <chr~
## 2 Dart~    175   80 none      red       yellow      54 male mascu~ Dathomir Zabrak <chr~
## 3 Ackb~    180   83 none      brown mot~ orange      41 male mascu~ Mon Cala Mon Ca~ <chr~
## 4 Luke~    172   77 blond fair      blue       19 male mascu~ Tatooine Human  <chr~
## 5 Zam ~    168   55 blonde fair, gre~ yellow      NA fema~ femin~ Zolan   Clawdi~ <chr~
## # ... with 2 more variables: vehicles <list>, starships <list>
```

## slice\_min() and slice\_max()

`slice_min()` and `slice_max()` lets you select rows with the lowest or highest values in a variable. It is similar to using `arrange()` on a single variable and then `head()`.

```
starwars %>% slice_max(mass, n = 3)
```

```
## # A tibble: 3 x 14
##   name    height  mass hair_color skin_color eye_color birth_year sex    gender homeworld species films
##   <chr>   <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr>   <chr>   <chr>   <lis>
## 1 Jabb~    175  1358 <NA>       green-tan~ orange        600 herm~ mascul~ Nal Hutta Hutt    <chr~
## 2 Grie~    216   159 none       brown, wh~ green, y~      NA male  mascul~ Kalee    Kaleesh <chr~
## 3 IG-88    200   140 none       metal      red          15 none  mascul~ <NA>     Droid   <chr~
## # ... with 2 more variables: vehicles <list>, starships <list>
```

# Create new variables with mutate()

Use `mutate()` to create new variables based on existing variables. The new variable will be the last column, so we frequently use it with `select`.

```
starwars %>%  
  mutate(height_in = height / 2.54) %>% head(1)
```

```
## # A tibble: 1 x 15  
##   name height mass hair_color skin_color eye_color birth_year sex gender homeworld species films  
##   <chr> <int> <dbl> <chr>      <chr>      <chr>      <dbl> <chr> <chr> <chr>      <chr> <lis>  
## 1 Luke~ 172 77 blond fair blue 19 male mascu~ Tatooine Human <chr~  
## # ... with 3 more variables: vehicles <list>, starships <list>, height_in <dbl>
```

```
starwars %>%  
  mutate(height_in = height / 2.54) %>%  
  select(name, height, height_in) %>% head(1)
```

```
## # A tibble: 1 x 3  
##   name          height height_in  
##   <chr>         <int>      <dbl>  
## 1 Luke Skywalker 172        67.7
```

# New variables must have the same number of rows

**Important:** Because `mutate()` adds a new column to the data set, the variable you are creating must have the same number of values as rows in the data set.

```
starwars %>%  
  select(name, mass) %>%  
  mutate(cumulative_mean = cummean(mass))
```

```
## # A tibble: 87 x 3  
##   name                mass cumulative_mean  
##   <chr>              <dbl>          <dbl>  
## 1 Luke Skywalker      77             77  
## 2 C-3P0               75             76  
## 3 R2-D2               32            61.3  
## 4 Darth Vader        136             80  
## 5 Leia Organa         49            73.8  
## 6 Owen Lars          120            81.5  
## 7 Beru Whitesun lars  75            80.6  
## 8 R5-D4               32            74.5  
## 9 Biggs Darklighter   84            75.6  
## 10 Obi-Wan Kenobi     77            75.7  
## # ... with 77 more rows
```

## Some useful functions for `mutate()`

- `pmin()`, `pmax()` Element-wise min and max
- `cummin()`, `cummax()` Cumulative min and max
- `cumsum()`, `cumprod()` Cumulative sum and product
- `between()` Are values between a and b?
- `cummean()` Cumulative mean
- `lead()`, `lag()` Copy values with offset
- `ntile()` Bin vector into n buckets

# mutate() examples

```
starwars %>%
  select(name, mass, birth_year) %>%
  mutate(
    cummin_mass = cummin(mass), # cummin gives the min value seen so far
    ratio = mass / mean(mass, na.rm = TRUE), # we divide mass/by the col mean
    massyear_pmin = pmin(mass, birth_year), # pmin gives the element-wise min
    lag2 = lag(massyear_pmin, 2)) # lag offsets the column values

## # A tibble: 87 x 7
##   name                mass birth_year cummin_mass ratio massyear_pmin lag2
##   <chr>              <dbl>    <dbl>      <dbl> <dbl>      <dbl> <dbl>
## 1 Luke Skywalker      77        19         77 0.791         19    NA
## 2 C-3PO               75       112         75 0.771         75    NA
## 3 R2-D2               32        33         32 0.329         32    19
## 4 Darth Vader        136       41.9         32 1.40         41.9   75
## 5 Leia Organa         49        19         32 0.504         19    32
## 6 Owen Lars          120        52         32 1.23         52   41.9
## 7 Beru Whitesun lars  75        47         32 0.771         47    19
## 8 R5-D4               32        NA         32 0.329         NA    52
## 9 Biggs Darklighter  84        24         32 0.863         24    47
## 10 Obi-Wan Kenobi     77        57         32 0.791         57    NA
```

## # ... with 77 more rows

Copyright Miles Chen. For personal use only. Do not distribute.

# Summarize with summarise()

Hadley is from New Zealand where they spell it with an s. He later added `summarize()` to have the same functionality, but I'm accustomed to using the original function.

Summary functions take multiple values and summarize them with a single value. For example, `mean()` and `var()` are summary functions.

```
starwars %>%
  select(height, mass) %>%
  summarise(
    avg_height = mean(height, na.rm = TRUE),
    var_height = var(height, na.rm = TRUE),
    avg_mass = mean(mass, na.rm = TRUE),
    min_height = min(height, na.rm = TRUE),
    max_mass = max(mass, na.rm = TRUE),
    count = n())
```

  

```
## # A tibble: 1 x 6
##   avg_height var_height avg_mass min_height max_mass count
##   <dbl>      <dbl>    <dbl>    <int>    <dbl> <int>
## 1      174.      1209.     97.3        66     1358    87
```

# Create groups using group\_by()

We can create groups using the group\_by() function.

```
starwars %>%  
  group_by(species) %>%  
  select(name, height, mass, species)
```

```
## # A tibble: 87 x 4  
## # Groups:   species [38]  
##   name          height  mass species  
##   <chr>         <int> <dbl> <chr>  
## 1 Luke Skywalker    172    77 Human  
## 2 C-3P0             167    75 Droid  
## 3 R2-D2              96    32 Droid  
## 4 Darth Vader       202   136 Human  
## 5 Leia Organa       150    49 Human  
## 6 Owen Lars         178   120 Human  
## 7 Beru Whitesun lars 165    75 Human  
## 8 R5-D4              97    32 Droid  
## 9 Biggs Darklighter 183    84 Human  
## 10 Obi-Wan Kenobi    182    77 Human  
## # ... with 77 more rows
```



# group\_by() + summarise()

The power of group\_by() is realized when combined with summarise()

```
starwars %>%  
  group_by(species) %>%  
  select(name, height, mass, species) %>%  
  summarise(  
    mean_ht = mean(height, na.rm = TRUE),  
    sd_ht = sd(height, na.rm = TRUE),  
    mean_mass = mean(mass, na.rm = TRUE),  
    sd_mass = sd(mass, na.rm = TRUE),  
    count = n())
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 38 x 6
```

```
##   species    mean_ht sd_ht mean_mass sd_mass count  
##   <chr>      <dbl> <dbl>      <dbl>  <dbl> <int>  
## 1 Aleena      79    NA        15      NA     1  
## 2 Besalisk   198    NA       102      NA     1  
## 3 Cerean     198    NA        82      NA     1  
## 4 Chagrian   196    NA       NaN      NA     1  
## 5 Clawdite   168    NA        55      NA     1  
## 6 Droid      131    49.1     69.8    51.0     6  
## 7 Dug       112    NA        40      NA     1
```

## group\_by() + summarise()

```
starwars %>%
  group_by(species) %>%
  select(name, height, mass, species) %>%
  summarise(
    mean_ht = mean(height, na.rm = TRUE),
    sd_ht = sd(height, na.rm = TRUE),
    mean_mass = mean(mass, na.rm = TRUE),
    sd_mass = sd(mass, na.rm = TRUE),
    count = n()
  ) %>%
  filter(count > 1) %>%
  arrange(desc(count)) %>%
  head()
```

## 'summarise()' ungrouping output (override with '.groups' argument)

## # A tibble: 6 x 6

##	species	mean_ht	sd_ht	mean_mass	sd_mass	count
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<int>
## 1	Human	177.	12.5	82.8	19.4	35
## 2	Droid	131.	49.1	69.8	51.0	6
## 3	<NA>	181.	2.89	48	NA	4
## 4	Gungan	209.	14.2	74	11.3	3

## group\_by() + mutate()

Note that C-3PO is above average when compared to other droids in the data set, but below average when compared to all characters in the data set.

```
starwars %>%  
  filter(species %in% c("Human", "Droid") |  
         is.na(species)) %>%  
  select(name, species, height) %>%  
  group_by(species) %>%  
  mutate(z_height = (height - mean(height, na.rm = TRUE))/sd(height, na.rm = TRUE)) %>%  
  head()
```

```
## # A tibble: 6 x 4  
## # Groups:   species [2]  
##   name          species height z_height  
##   <chr>         <chr>    <int>    <dbl>  
## 1 Luke Skywalker Human      172    -0.371  
## 2 C-3PO         Droid      167     0.728  
## 3 R2-D2         Droid       96    -0.716  
## 4 Darth Vader   Human      202     2.02  
## 5 Leia Organa   Human      150    -2.13  
## 6 Owen Lars     Human      178     0.108
```

## Without group\_by()

Note that C-3PO is above average when compared to other droids in the data set, but below average when compared to all characters in the data set.

```
starwars %>%  
  filter(species %in% c("Human", "Droid") |  
         is.na(species)) %>%  
  select(name, species, height) %>%  
  # group_by(species) %>%  
  mutate(z_height = (height - mean(height, na.rm = TRUE))/sd(height, na.rm = TRUE)) %>%  
  head()
```

```
## # A tibble: 6 x 4  
##   name      species height z_height  
##   <chr>      <chr>    <int>   <dbl>  
## 1 Luke Skywalker Human      172    0.0329  
## 2 C-3PO      Droid      167   -0.168  
## 3 R2-D2      Droid       96   -3.02  
## 4 Darth Vader Human      202    1.24  
## 5 Leia Organa Human      150   -0.849  
## 6 Owen Lars  Human      178    0.274
```

# Multiple group\_by() on some toy data

```
toy_cases <- read_csv("https://raw.githubusercontent.com/rstudio/EDAWR/master/data-raw/toyb.csv")
print(toy_cases)
```

```
## # A tibble: 12 x 4
##   country      year sex    cases
##   <chr>      <dbl> <chr> <dbl>
## 1 Afghanistan 1999 female     1
## 2 Afghanistan 1999 male      1
## 3 Afghanistan 2000 female     1
## 4 Afghanistan 2000 male      1
## 5 Brazil      1999 female     2
## 6 Brazil      1999 male      2
## 7 Brazil      2000 female     2
## 8 Brazil      2000 male      2
## 9 China       1999 female     3
## 10 China      1999 male      3
## 11 China      2000 female     3
## 12 China      2000 male      3
```

# Multiple group\_by() + summarise()

We can provide group\_by() two variables and it will create a hierarchy of groups

```
summary1 <- toy_cases %>% group_by(country, year) %>%  
  summarise(cases = sum(cases))
```

```
## 'summarise()' regrouping output by 'country' (override with '.groups' argument)
```

```
print(summary1)
```

```
## # A tibble: 6 x 3  
## # Groups:   country [3]  
##   country      year cases  
##   <chr>      <dbl> <dbl>  
## 1 Afghanistan 1999     2  
## 2 Afghanistan 2000     2  
## 3 Brazil       1999     4  
## 4 Brazil       2000     4  
## 5 China        1999     6  
## 6 China        2000     6
```

# Multiple group\_by() + summarise()

```
summary2 <- summary1 %>% summarise(cases = sum(cases))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
summary2
```

```
## # A tibble: 3 x 2
##   country    cases
##   <chr>      <dbl>
## 1 Afghanistan     4
## 2 Brazil           8
## 3 China          12
```

```
summary3 <- summary2 %>% summarise(cases = sum(cases))
```

```
summary3
```

```
## # A tibble: 1 x 1
##   cases
##   <dbl>
## 1     24
```

# Change the order of the multiple group\_by()

```
summary_a <- toy_cases %>% group_by(year, country) %>%  
  summarise(cases = sum(cases))
```

```
## 'summarise()' regrouping output by 'year' (override with '.groups' argument)
```

```
print(summary_a)
```

```
## # A tibble: 6 x 3  
## # Groups:   year [2]  
##   year country    cases  
##   <dbl> <chr>      <dbl>  
## 1  1999 Afghanistan    2  
## 2  1999 Brazil        4  
## 3  1999 China         6  
## 4  2000 Afghanistan    2  
## 5  2000 Brazil        4  
## 6  2000 China         6
```



# Multiple group\_by() + summarise()

```
summary_b <- summary_a %>% summarise(cases = sum(cases))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
summary_b
```

```
## # A tibble: 2 x 2
```

```
##   year cases
```

```
##   <dbl> <dbl>
```

```
## 1  1999     12
```

```
## 2  2000     12
```

```
summary_c <- summary_b %>% summarise(cases = sum(cases))
```

```
summary_c
```

```
## # A tibble: 1 x 1
```

```
##   cases
```

```
##   <dbl>
```

```
## 1     24
```

## Section 3

### Two-table verbs

# Two-table verbs

dplyr also comes with some two-table verbs that allow you to combine tables.

**Mutating joins** add new variables to one table from matching rows in another

Not covered here, but you can read more at <https://dplyr.tidyverse.org/articles/two-table.html>

- **Filtering joins** filter observations from one table based on whether or not they match an observation in the other.
- **Set operations**, which combine the observations in the data sets as if they were set elements.

# Toy tables

```
people <- tibble(  
  name = c("Adam", "Betty", "Carl", "Doug"),  
  state = c("CA", "CA", "NY", "TX")  
)  
  
states <- tibble(  
  abbreviation = c("CA", "NY", "WA"),  
  state_name = c("California", "New York", "Washington")  
)
```

# left\_join()

`left_join()` takes all the values in the left table and adds variables from the right table by matching values using a column that exists in both tables. Values that do not exist in the other table have NA returned.

```
people %>% left_join(states, by = c("state" = "abbreviation"))
```

```
## # A tibble: 4 x 3
##   name  state state_name
##   <chr> <chr> <chr>
## 1 Adam  CA      California
## 2 Betty CA      California
## 3 Carl  NY      New York
## 4 Doug  TX      <NA>
```

# right\_join()

`right_join()` is similar to `left_join` except it keeps all the rows in the right table.

```
people %>% right_join(states, by = c("state" = "abbreviation"))
```

```
## # A tibble: 4 x 3
##   name  state state_name
##   <chr> <chr> <chr>
## 1 Adam  CA      California
## 2 Betty CA      California
## 3 Carl  NY      New York
## 4 <NA>  WA      Washington
```

# inner\_join()

`inner_join()` keeps only rows that have values that exist in both tables. You can think of this as the intersection.

```
people %>% inner_join(states, by = c("state" = "abbreviation"))
```

```
## # A tibble: 3 x 3
##   name  state state_name
##   <chr> <chr> <chr>
## 1 Adam  CA      California
## 2 Betty CA      California
## 3 Carl  NY      New York
```

# full\_join()

`full_join()` keeps all rows from both tables. You can think of this as the union. (in SQL this is called a full outer join)

```
people %>% full_join(states, by = c("state" = "abbreviation"))
```

```
## # A tibble: 5 x 3
##   name  state state_name
##   <chr> <chr> <chr>
## 1 Adam  CA      California
## 2 Betty CA      California
## 3 Carl  NY      New York
## 4 Doug  TX      <NA>
## 5 <NA>  WA      Washington
```



# Controlling how the tables are matched

Depending on the tables, the join operation can match tables on different variables.

In the previous examples, we used a named character vector `by = c("state" = "abbreviation")` specifying the name in the left table that matches the name in the right table.

## Options for joining tables

- `by = NULL` (or don't specify anything): `dplyr` will use all variables that have the same name in both tables.
- `by = "x"`: `dplyr` will use only some of the variables that have the same name in both tables
- `by = c("x" = "y")`: this is the form that must be used if the matching columns do not have the same name in both tables.